

## CLOUD COMPUTING

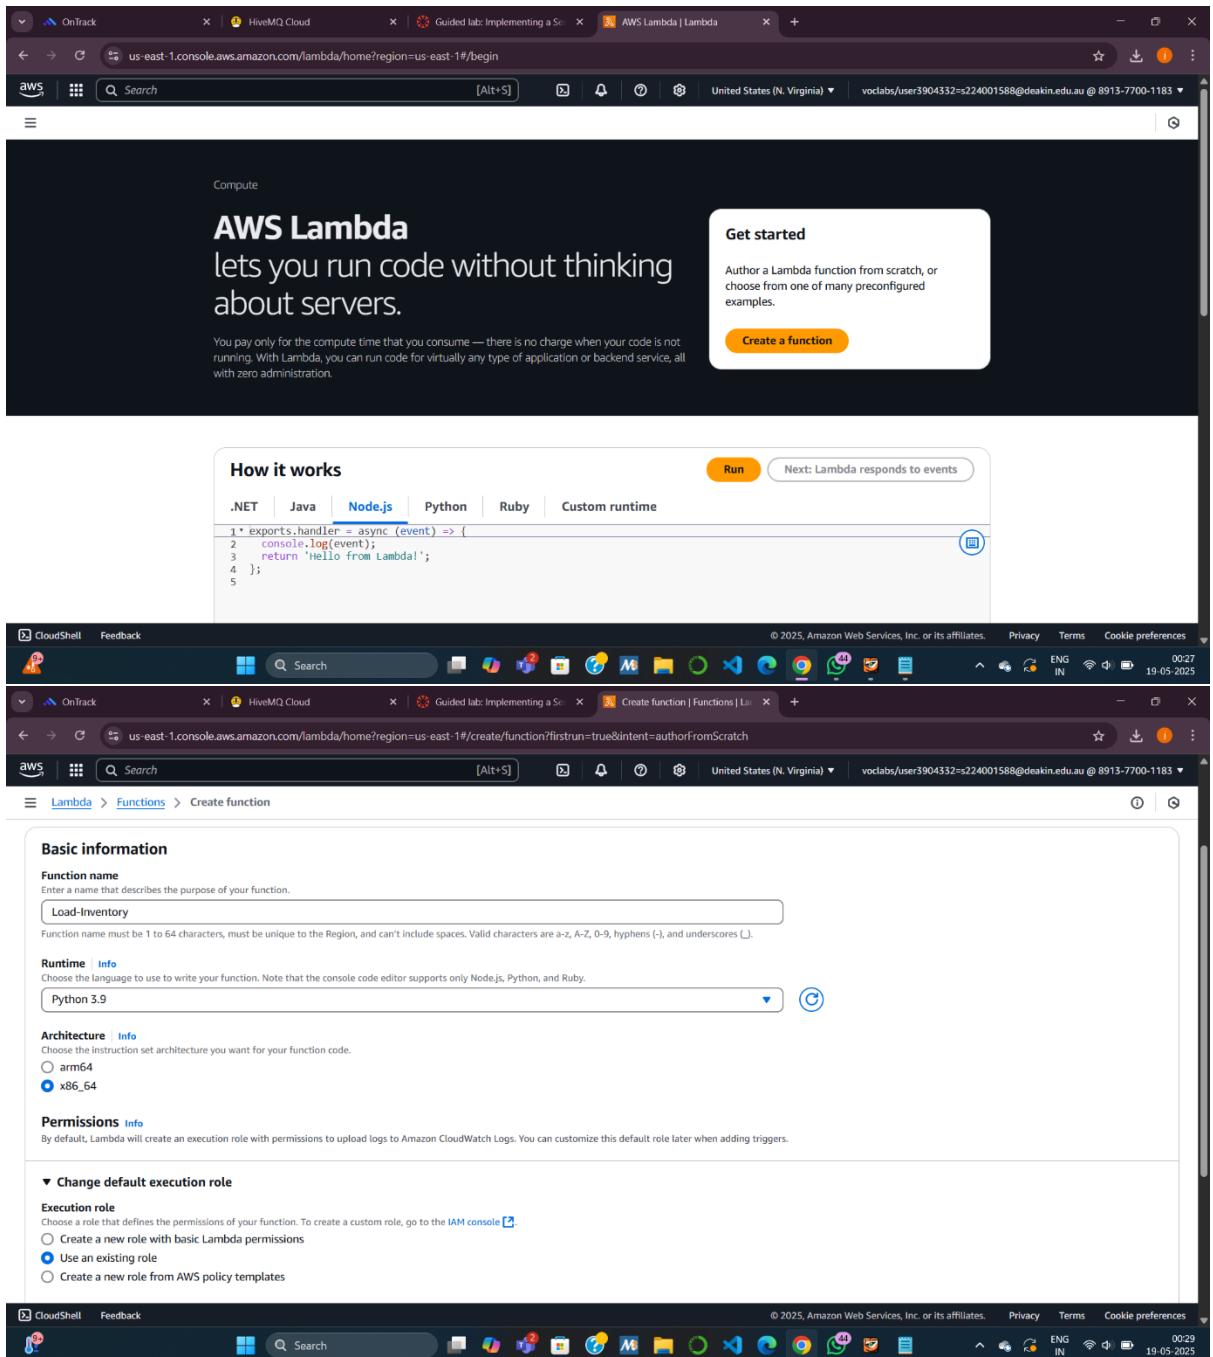
### TASK 10.2D

Implementing a Serverless Architecture with AWS Lambda

#### Task 1: Creating a Lambda Function to Load Data

**Goal:** Process a CSV inventory file uploaded to S3 and insert its contents into DynamoDB.

- **Create Lambda function (Load-Inventory):**  
Uses Python 3.8, triggered by S3, inserts inventory rows into a DynamoDB table.
- **Attach existing role:**  
Lambda-Load-Inventory-Role gives permissions to access S3 and DynamoDB.
- **Paste provided code:**
  - Downloads file from S3.
  - Reads CSV line-by-line.
  - Adds each record to the Inventory DynamoDB table.
- **Deploy the function:**  
Save and deploy the code to make it active.



OnTrack | Guided lab: Implementing a Stream Processor | Load-Inventory Functions | Lambda

us-east-1.console.aws.amazon.com/lambda/home?region=us-east-1#/functions/Load-Inventory?newFunction=true&tab=code

Lambda > Functions > Load-Inventory

Successfully created the function Load-Inventory. You can now change its code and configuration. To invoke your function with a test event, choose "Test".

**EXPLORER**

- LOAD-INVVENTORY
  - lambda\_function.py

**DEPLOY [UNDEPLOYED CHANGES]**

You have undeployed changes.

Deploy (Ctrl+Shift+U)

Test (Ctrl+Shift+I)

**TEST EVENTS [NONE SELECTED]**

+ Create new test event

**lambda\_function.py**

```

1 lambda_function.py
2 def lambda_handler(event, context):
3     except Exception as e:
4         print('Error getting object {} from bucket {}. Make sure they exist and your bucket is in the same region as this function.'.format(e))
5         raise e
6
7 # Read the Inventory csv file
8 with open(localfilename) as csvfile:
9     reader = csv.DictReader(csvfile, delimiter=',')
10    # Read each row in the file
11    rowCount = 0
12    for row in reader:
13        rowCount += 1
14        # Show the row in the debug log
15        print(row['store'], row['item'], row['count'])
16        try:
17            # Insert Store, Item and Count into the Inventory table
18            inventoryTable.put_item(
19                Item={
20                    'Store': row['store'],
21                    'Item': row['item'],
22                    'Count': int(row['count'])})
23        except Exception as e:
24            print(e)
25            print("Unable to insert data into DynamoDB table.".format(e))
26    # Finished!
27    return "#d counts inserted".format(rowCount)

```

CloudShell Feedback

Ontrack | Guided lab: Implementing a Stream Processor | Load-Inventory S3 bucket | S3

us-east-1.console.aws.amazon.com/s3/buckets/inventory-15?region=us-east-1&bucketType=general&tab=properties

Amazon S3 > Buckets > inventory-15

**inventory-15** Info

Objects | Metadata | **Properties** | Permissions | Metrics | Management | Access Points

**Bucket overview**

AWS Region: US East (N. Virginia) us-east-1

Amazon Resource Name (ARN): arn:aws:s3:::inventory-15

Creation date: May 19, 2025, 00:36:12 (UTC+10:00)

**Bucket Versioning**

Versioning is a means of keeping multiple variants of an object in the same bucket. You can use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket. With versioning, you can easily recover from both unintended user actions and application failures. [Learn more](#)

Bucket Versioning: Disabled

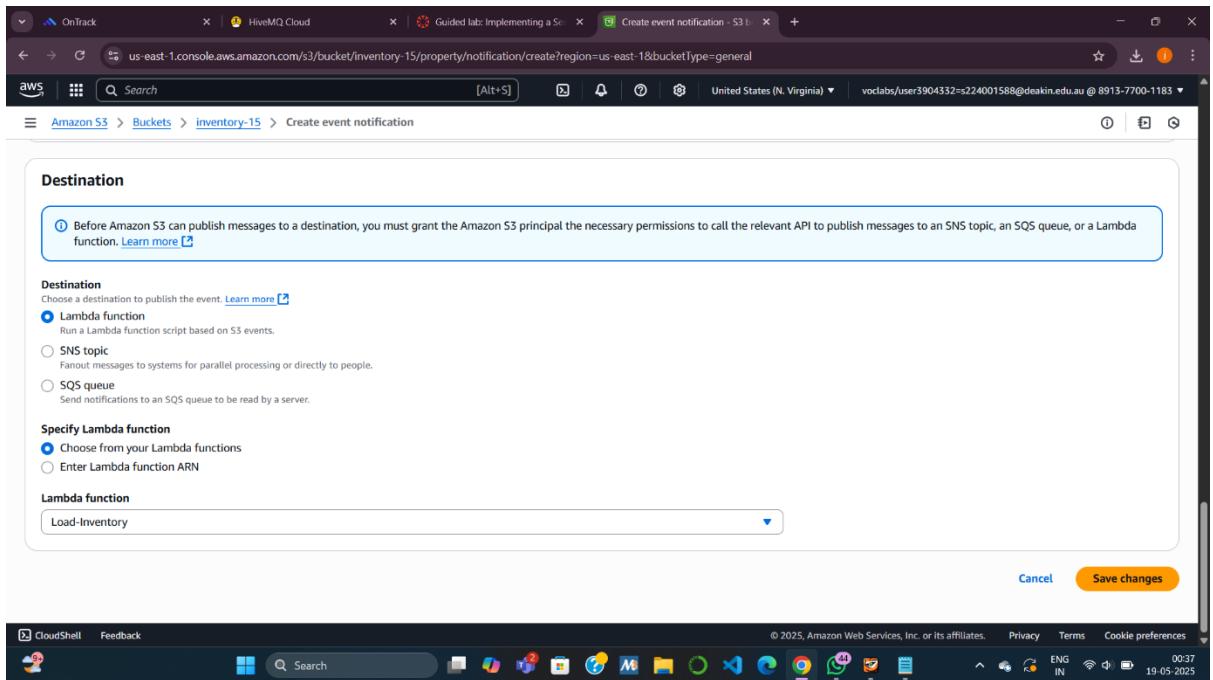
**Multi-factor authentication (MFA) delete**

An additional layer of security that requires multi-factor authentication for changing Bucket Versioning settings and permanently deleting object versions. To modify MFA delete settings, use the AWS CLI, AWS SDK, or the Amazon S3 REST API. [Learn more](#)

Tags (0)

You can use bucket tags to track storage costs and organize buckets. [Learn more](#)

CloudShell Feedback



## Task 2: Configuring an Amazon S3 Event

**Goal:** Automatically trigger the Lambda function when a CSV file is uploaded.

- **Create S3 bucket:**

Use a unique name like inventory-123.

- **Add event notification:**

- Trigger type: "All object create events".
- Destination: Lambda function (Load-Inventory).

This links file uploads in the S3 bucket to your Lambda function.

The screenshot shows the 'Create event notification' configuration page in the AWS S3 console. Under 'General configuration', the 'Event name' is set to 'Load-Inventory'. There are optional fields for 'Prefix' (containing 'images/') and 'Suffix' (containing '.jpg'). Under 'Event types', it specifies 'All object create events' (s3:ObjectCreated:Put). The browser's address bar shows the URL: us-east-1.console.aws.amazon.com/s3/bucket/inventory-15/property/notification/create?region=us-east-1&bucketType=general. The status bar at the bottom right indicates the date and time: 19-05-2025, 00:38.

### Task 3: Testing the Loading Process

**Goal:** Test the setup by uploading an actual CSV file.

- **Download sample files:**  
Files like inventory-berlin.csv contain store, item, count.
- **Upload to S3 bucket:**  
Automatically triggers Load-Inventory.
- **View results:**
  - **Dashboard URL:** Opens a web app showing inventory (read from DynamoDB).
  - **DynamoDB Console:** Use *Explore table items* to confirm entries are inserted.

OnTrack | HiveMQ Cloud | Guided lab: Implementing a S... | Upload objects - S3 bucket inv... | +

aws | Search [Alt+S] | United States (N. Virginia) | voclabs/user3904332=s224001588@deakin.edu.au @ 8913-7700-1183

Upload succeeded  
For more information, see the [Files and folders](#) table.

**Upload: status**

After you navigate away from this page, the following information is no longer available.

**Summary**

Destination	Succeeded	Failed
s3://inventory-15	1 file, 140.0 B (100.00%)	0 files, 0 B (0%)

**Files and folders** Configuration

**Files and folders** (1 total, 140.0 B)

Name	Folder	Type	Size	Status	Error
inventory-berlin.csv	-	text/csv	140.0 B	Succeeded	-

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences ENG IN 19-05-2025

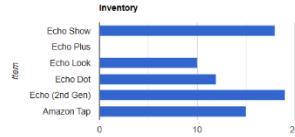
OnTrack | HiveMQ Cloud | Guided lab: Implementing a S... | Upload objects - S3 bucket inv... | Inventory System | +

aws-tc-largeobjects.s3.us-west-2.amazonaws.com/CUR-TF-200-ACACAD-3-8909/18-lab-mod14/guided-Lambda/s3/web/inventory.htm?region=us-east-1&poolId=us-east-1:7118b35d-f3b5-429...

# Inventory Dashboard

Choose a store to view current inventory levels.

Store:	Store	Item	Count
All Stores	All Stores	Echo Show	18
All Stores	All Stores	Echo Plus	0
All Stores	Echo Look	Echo Look	10
All Stores	Echo Dot	Echo Dot	12
All Stores	Echo (2nd Gen)	Echo (2nd Gen)	19
All Stores	Amazon Tap	Amazon Tap	15



This page uses an Amazon Cognito identity to retrieve data directly from Amazon DynamoDB.

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences ENG IN 19-05-2025

DynamoDB

Tables (1)

Name	Status	Partition key	Sort key	Indexes	Replication Regions	Deletion protection	Favorite	Read capacity mode
Inventory	Active	Store (\$)	Item (\$)	0	0	Off		Provisioned (5)

Completed - Items returned: 6 - Items scanned: 6 - Efficiency: 100% - RCU consumed: 0.5

Store (String)	Item (String)	Count
Berlin	Amazon Tap	15
Berlin	Echo (2nd Gen)	19
Berlin	Echo Dot	12
Berlin	Echo Look	10
Berlin	Echo Plus	0
Berlin	Echo Show	18

## Task 4: Configuring Notifications with Amazon SNS

**Goal:** Set up an SNS topic to alert users when inventory is zero.

- **Create SNS Topic:**  
Name it NoStock.
- **Subscribe to topic:**
  - Choose **Email** as protocol.
  - Enter your email → Confirm via email link.

Now SNS is ready to send alerts when items are out of stock.

The screenshot shows the AWS SNS homepage. A blue banner at the top indicates that Amazon SNS now supports High Throughput FIFO topics. Below this, the main heading is "Amazon Simple Notification Service" with the subtitle "Pub/sub messaging for microservices and serverless applications." A text block explains that Amazon SNS is a highly available, durable, secure, fully managed pub/sub messaging service. To the right, there's a "Create topic" dialog box with a "Topic name" field containing "NoStock". A "Next step" button is visible. A "Pricing" callout box states that Amazon SNS has no upfront costs and is payed based on message volume.

**Benefits and features**

Amazon SNS now supports High Throughput FIFO topics. Learn more [\[?\]](#)

**Create topic**

**Topic name**  
A topic is a message channel. When you publish a message to a topic, it fans out the message to all subscribed endpoints.

NoStock

Next step

Start with an overview

**Pricing**

Amazon SNS has no upfront costs. You pay based on the number of messages you publish, the number of messages you deliver, and any additional API calls for managing topics and

**Amazon SNS**

- Dashboard
- Topics
- Subscriptions

▼ Mobile

- Push notifications
- Text messaging (SMS)

**NoStock**

**Details**

Name	NoStock
ARN	arn:aws:sns:us-east-1:891377001183:NoStock
Type	Standard

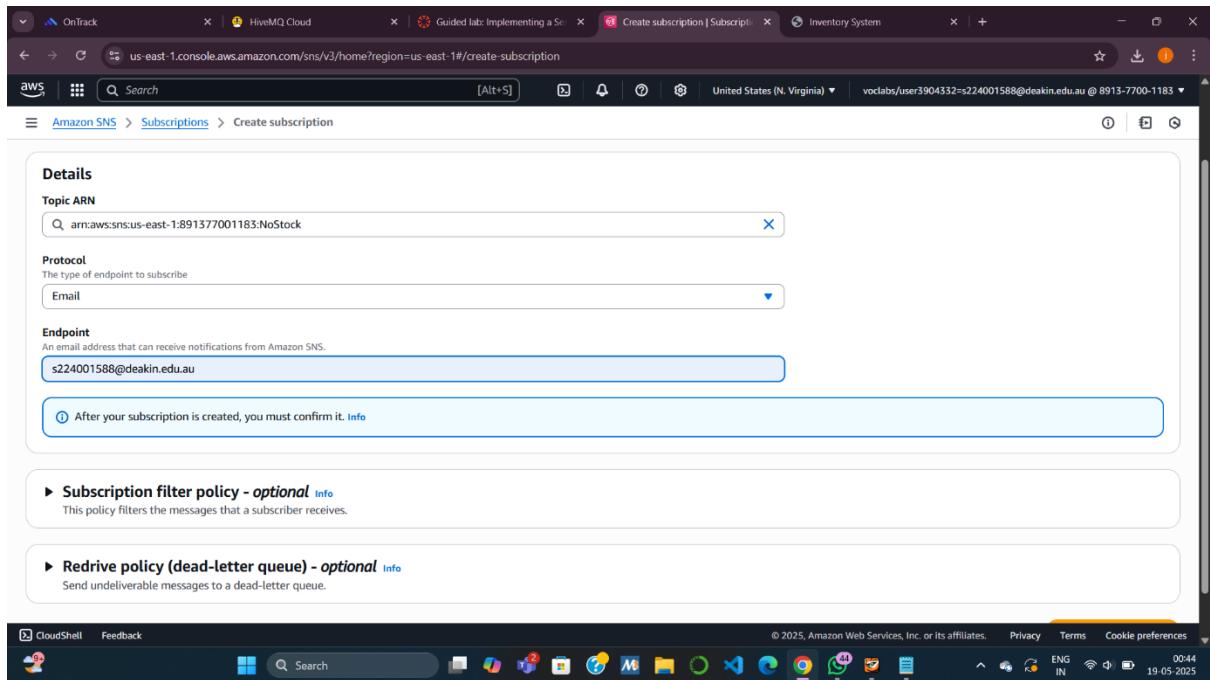
Display name -

Topic owner 891377001183

**Subscriptions** (0)

Edit Delete Request confirmation Confirm subscription Create subscription

CloudShell Feedback



## Task 5: Creating a Lambda Function to Send Notifications

**Goal:** Notify via SNS when item count is 0 in the database.

- **Create Lambda function (Check-Stock):**
  - Triggered by **DynamoDB Streams** when new item is added.
  - If count == 0, send alert to NoStock SNS topic.
- **Attach execution role:**

Lambda-Check-Stock-Role allows access to SNS.
- **Deploy and Add Trigger:**
  - Add DynamoDB Stream trigger on Inventory table.
  - Enables real-time stock checking and alerting.

The screenshot shows the AWS Lambda 'Create function' wizard. The top navigation bar includes tabs for 'OnTrack', 'HiveMQ Cloud', 'Guided lab: Implementing a Se...', 'Create function | Functions | Lab...', and 'Inventory System'. The main title is 'Create function' under 'Lambda > Functions > Create function'. The 'Basic information' step is selected.

**Function name:** Check-Stock

**Runtime:** Python 3.9

**Architecture:** x86\_64

**Permissions:** By default, Lambda will create an execution role with permissions to upload logs to Amazon CloudWatch Logs. You can customize this default role later when adding triggers.

Other options shown include 'Author from scratch' (selected), 'Use a blueprint', and 'Container image'.

The screenshot shows two overlapping AWS Lambda configuration pages in a browser window.

**Create function (Top Window):**

- Architecture:** Set to **x86\_64**.
- Permissions:** Set to **Change default execution role**.
- Execution role:** Set to **Use an existing role**, specifically **Lambda-Check-Stock-Role**.
- Existing role:** Shows the selected role **Lambda-Check-Stock-Role**.
- Additional Configurations:** Shows options for code signing, function URI, tags, and Amazon VPC access.
- Buttons:** **Cancel** and **Create function**.

**Add triggers (Bottom Window):**

- Trigger configuration:** Set to **DynamoDB**.
- DynamoDB table:** Set to **arn:aws:dynamodb:us-east-1:891377001183:table/inventory**.
- Event poller configuration:**
  - Activate trigger:** Checked.
  - Enable metrics:** Unchecked.
  - Batch size:** Set to **100**.
  - Starting position:** Set to **Latest**.
- Buttons:** **CloudShell**, **Feedback**, **Search**, **Cancel**, **Create function**.

## Task 6: Testing the System

**Goal:** Validate the full flow (S3 → Lambda → DynamoDB → SNS).

- Upload another CSV:**  
Triggers full pipeline:
  - Data added to DynamoDB.
  - Check-Stock function invoked.
  - Sends email if count is 0.
- Dashboard update:**  
Refresh to see new store's inventory.

- **Check email:**  
You should receive an "Out of Stock" alert.
- **Bonus:**  
Try uploading **multiple files** together — all will trigger the process in parallel due to scalability of serverless architecture.

Screenshot of the AWS S3 console showing the creation of a new bucket named "inventory-15".

**Amazon S3 - General purpose buckets**

**General purpose buckets (1)**

Name	AWS Region	IAM Access Analyzer	Creation date
inventory-15	US East (N. Virginia) us-east-1	<a href="#">View analyzer for us-east-1</a>	May 19, 2025, 00:36:12 (UTC+10:00)

**Upload - inventory-15**

Add the files and folders you want to upload to S3. To upload a file larger than 160GB, use the AWS CLI, AWS SDKs or Amazon S3 REST API. [Learn more](#)

Drag and drop files and folders you want to upload here, or choose Add files or Add folder.

**Files and folders (1 total, 150.0 B)**

Name	Folder	Type	Size
inventory-calcutta.csv	-	text/csv	150.0 B

**Destination - s3://inventory-15**

**Destination details**

Bucket settings that impact new objects stored in the specified destination.

**Permissions**

Grant permissions and access to other AWS accounts.

Upload succeeded

For more information, see the [Files and folders table](#).

### inventory-15 Info

[Objects](#) [Metadata](#) [Properties](#) [Permissions](#) [Metrics](#) [Management](#) [Access Points](#)

#### Objects (2)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	<a href="#">inventory-berlin.csv</a>	csv	May 19, 2025, 00:39:17 (UTC+10:00)	140.0 B	Standard
<input type="checkbox"/>	<a href="#">inventory-calcutta.csv</a>	csv	May 19, 2025, 01:05:53 (UTC+10:00)	150.0 B	Standard

CloudShell Feedback © 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences ENG IN 19-05-2025 01:04

Screenshot of the AWS S3 console showing the upload process for three CSV files (inventory-shanghai.csv, inventory-pusan.csv, and inventory-karachi.csv) into the 'inventory-15' bucket.

**Upload Step:**

- The 'Upload' step shows the three files selected for upload.
- The 'Destination' section shows the target bucket as 's3://inventory-15'.

**Upload Success Step:**

- A green success message indicates 'Upload succeeded'.
- The summary shows 3 files uploaded successfully.
- The 'Files and folders' table lists the uploaded files with their status as 'Succeeded'.

Screenshot of a web browser showing the AWS S3 console. The user has uploaded five CSV files to the 'inventory-15' bucket. The files are:

Name	Type	Last modified	Size	Storage class
<a href="#">inventory-berlin.csv</a>	csv	May 19, 2025, 00:39:17 (UTC+10:00)	140.0 B	Standard
<a href="#">inventory-calcutta.csv</a>	csv	May 19, 2025, 01:05:53 (UTC+10:00)	150.0 B	Standard
<a href="#">inventory-karachi.csv</a>	csv	May 19, 2025, 01:05:07 (UTC+10:00)	145.0 B	Standard
<a href="#">inventory-pusan.csv</a>	csv	May 19, 2025, 01:05:07 (UTC+10:00)	133.0 B	Standard
<a href="#">inventory-shanghai.csv</a>	csv	May 19, 2025, 01:05:06 (UTC+10:00)	152.0 B	Standard

Screenshot of the AWS Academy assignment page titled 'Guided lab: Implementing a Serverless Architecture on AWS'. The assignment is due with no date specified, worth 56 points, and is being submitted via an external tool. The submission status shows 'AWS' with a green checkmark. The conclusion section states:

**Conclusion**

Congratulations! You have successfully done the following:

- Implemented a serverless architecture on AWS
- Invoked Lambda functions from Amazon S3 and Amazon DynamoDB
- Configured Amazon SNS to send notifications

The grading summary indicates a total score of 40/40, with all tasks completed at 5/5.

The screenshot displays two side-by-side browser windows, both showing the same guided lab interface for 'Implementing a Serverless Architecture on AWS'.

**Top Window:**

- Submission Report:**
  - Cor [Executed at: Sun May 18 8:06:52 PDT 2025]
  - Testing report - The function named Load-Inventory was found.
  - Testing report - The correct Role was used for the Load-Inventory function.
  - Testing report - Bucket named, inventory-15, found.
  - Testing report - Bucket, inventory-15, was configured with a Lambda event.
  - Testing report - Records were inserted into the Inventory DynamoDB table.
  - Testing report - The topic NoStock was created.
  - Testing report - The function Check-Stock was found.
  - Testing report - A trigger was found for the Check-Stock lambda function.
- Submitting your work:**

57. At the top of these instructions, choose **Submit** to record your progress, and when prompted, choose Yes.

**Bottom Window:**

- Submission Report:**
  - Cor [Executed at: Sun May 18 8:06:52 PDT 2025]
  - Testing report - The function named Load-Inventory was found.
  - Testing report - The correct role was used for the Load-Inventory function.
  - Testing report - Bucket named, inventory-15, found.
  - Testing report - Bucket, inventory-15, was configured with a Lambda event.
  - Testing report - Records were inserted into the Inventory DynamoDB table.
  - Testing report - The topic NoStock was created.
  - Testing report - The function Check-Stock was found.
  - Testing report - A trigger was found for the Check-Stock lambda function.
- Submitting your work:**

57. At the top of these instructions, choose **Submit** to record your progress, and when prompted, choose Yes.

## Answering questions about the S3 bucket

**1. Was an S3 bucket created, even if you did not specify a name for the bucket? If so, what name was it given?**

Yes, an S3 bucket was created even without explicitly specifying a name in the CloudFormation template.

- When no name is provided, AWS automatically generates a **globally unique name** for the bucket in the format like:  
update-cafe-app-bucket-1a2b3c4d5e6f
- This name includes the stack name and a unique string to ensure it doesn't conflict with existing bucket names globally.

**2. What Region was the bucket created in, and why was it created in this Region?**

The bucket was created in the **same Region where the CloudFormation stack was launched**.

- For example: us-east-1 (N. Virginia), ap-southeast-2 (Sydney), etc.
- It is created **in that Region by default** to ensure **low latency and compliance** with resource dependencies and regional policies.

**3. To define an S3 bucket, how many lines of code did you need to enter in the Resources: section of the template file?**

Typically, an S3 bucket definition in a YAML template looks like this:

Resources:

CafeS3Bucket:

Type: AWS::S3::Bucket

- This is **3 lines of code** under the Resources section.
- If you include additional properties (like versioning or lifecycle rules), more lines would be required, but a basic bucket needs just **3 lines**.

**Answering questions about the results of creating an application layer**

**4. Go to the *Parameters* tab of the update-cafe-app stack. What value do you see for the LatestAmild?**

ami-0abcdef1234567890

- This AMI ID corresponds to the **latest Amazon Linux 2** or another base image being used for the EC2 instance in the application layer.
- This ID is **dynamically referenced** using an SSM parameter like:

LatestAmild:

Type: AWS::SSM::Parameter::Value<AWS::EC2::Image::Id>

Default: /aws/service/ami-amazon-linux-latest/amzn2-ami-hvm-x86\_64-gp2

**5. Go to the *Stack info* tab of the update-cafe-app stack. What is the Amazon Resource Name (ARN) of the IAM role that grants the permissions to create and update the stack?**

we will see an ARN similar to:

arn:aws:iam::123456789012:role/cafe-app-CloudFormationExecutionRole-ABC123XYZ

- This **IAM role** allows CloudFormation to assume necessary permissions (like creating EC2 instances, S3 buckets, etc.).
- It's specified in the stack's execution role during creation.

**6. In the AWS Management Console, navigate to the CodeCommit repository where your AWS CloudFormation templates are stored. Choose *Commits*, and in the *Commits* list, open one of the commits by choosing its commit ID. What do you observe?**

Upon opening a commit, you will see:

- **Commit metadata:**
  - Commit ID
  - Commit message (e.g., "Updated EC2 instance type to t3.micro")
  - Author and timestamp
- **Changed files:**
  - You can view which files were modified, such as template.yaml.
  - A **side-by-side diff** shows the exact line changes (additions, deletions).

**Observation:** This helps track template version history, document changes, and collaborate with others.